

Supplementary Material for Correct Predictions, Wrong Reasons: On the Limits of Metrics for Measuring Reasoning Ability

ANONYMOUS AUTHOR(S)

This supplementary material documents source separation, formulas, secondary-analysis protocols, complete constraint tables, robustness summaries, and architecture-level instantiations used in the main paper. The matched and shortcut-account analyses use values printed to three decimals in frozen source exports; the group analysis uses full-precision raw prediction records. The material deliberately treats conventional, matched, confidence-bounded, and shortcut-adjusted quantities as observable behavioral criteria rather than as direct measures of a latent reasoning mechanism.

Additional Key Words and Phrases: reasoning evaluation, metric validity, matched evaluation, witnesses, shortcut learning

Organization. Section S1 documents provenance and version separation. Section S2 defines the secondary analyses, including the full-precision group outputs. Sections S3–S5 report aggregate results, sensitivity analyses, and complete constraint tables. Section S6 provides architecture-level instantiations of the main argument, and Section S7 records reproduction requirements and remaining evidence boundaries.

S1 Provenance, Versioning, and Claim Discipline

The analysis uses two frozen constraint-evaluation archives and a separate full-precision group analysis. We refer to the frozen sources as the *matched-separability archive* and the *shortcut-margin archive*. The matched-separability archive is the sole source for row-level metric, witness, confidence-bound, and pair-rematching analyses. The shortcut-margin archive is used only for its ideal-floor versus residual-shortcut-ceiling boundary analysis and its point-estimate versus confidence-bound examples.

The group analysis reads current raw prediction records separately for 3-SAT, 2-SAT, and graph coloring. A single domain-neutral implementation then computes equal-group, equal-pair, nonlinear, variance, and repeated-success summaries. These outputs are not numerically merged with rows from either frozen archive.

This separation is necessary because some overlapping narrative rows differ slightly across the two archives, consistent with reruns, parser changes, or versioned exports. We do not average or selectively reconcile those values. A submission artifact should replace the source-export identifiers in its manifest with immutable repository commit identifiers and include the full-precision per-instance predictions.

The supplied constraint tables contain 19 3-SAT rows, 59 2-SAT rows, and 45 graph-coloring rows, for 123 model-size observations. The group analysis contains 19, 66, and 45 model-size slices for the same domains, respectively, for 130 slices in total.

S2 Secondary-Analysis Definitions

S2.1 Witness Reference and Two-Sided Separability

For SAT, the witness reference is ADR^{+w} : a positive response receives credit only when the emitted assignment satisfies the formula. For graph coloring, the analogous quantity is denoted ADR^{+a} : the

proposed coloring must be valid. In the frozen matched-separability table, both appear in the same source column; this supplement writes the task-appropriate symbol W .

For each matched group G , let $\hat{\theta}_+(G)$ and $\hat{\theta}_-(G)$ be positive- and negative-arm recall. The point two-sided separability is

$$\hat{s}_G = \hat{\theta}_+(G)\hat{\theta}_-(G).$$

The source archive reports the group-averaged point estimate together with a factorized Clopper–Pearson lower bound L_{CP} and Wilson lower bound L_W . Because $\hat{s}_G = G\text{-mean}^2$, the point statistic induces the same classifier ordering as $G\text{-mean}$. The contribution of the source protocol is the matched-group and confidence construction, not a new scalar ordering.

S2.2 Strict Ordering Discordance

Within each task and problem size, consider model pairs for which both a metric M and W differ at the reported precision. We compute

$$D_M = \frac{\sum_N \sum_{a < b} \mathbf{1}[(M_{a,N} - M_{b,N})(W_{a,N} - W_{b,N}) < 0]}{\sum_N \sum_{a < b} \mathbf{1}[M_{a,N} \neq M_{b,N} \wedge W_{a,N} \neq W_{b,N}]}.$$

No hidden precision is used to break ties. The statistic measures agreement with a stronger observable reference; it does not assume that W is a complete measure of reasoning.

S2.3 Metric-Level Gap

For each row and metric M , we record whether $M > W$, $M = W$, or $M < W$ at source precision and report the mean difference $M - W$. The scales have different operational meanings, so the difference is descriptive rather than an estimate of “overstatement” in a common latent unit.

S2.4 Pair-Rematching Sensitivity

The source tables include the original pairing, a maximum-rematching branch, and a minimum-rematching branch. We compute

$$\Delta_{\text{match}} = \text{ADR}_{\text{max}} - \text{ADR}_{\text{min}}.$$

This quantity changes only the matching rule applied to a fixed set of predictions.

S2.5 Confidence and Boundary Sensitivity

We summarize

$$\Delta_{CI} = |L_{CP} - L_W|$$

for every constraint row. Threshold-parametric prefix boundaries are read as

$$N^*(\tau) = \max \{N : L_{CP}(n) \geq \tau \text{ for every tested } n \leq N\},$$

subject to the source archive’s prefix convention.

The shortcut-margin archive additionally defines an ideal floor $C_0 = 0.25$ and a measured residual shortcut ceiling. Its pre-registered residual feature set is

$$\mathcal{F}_{CB} = \{\text{mean_clause_len}, \text{min_clause_len}, \text{max_clause_len}, \text{pos_ratio}, \text{var_coverage}\}.$$

For each group, the archive takes the maximum two-sample KS distance over these one-dimensional features, subtracts the 0.95 quantile of a within-group label-permutation reference computed from 200 permutations, and truncates at zero. Calling the resulting estimate $\hat{d}_{\text{short}}$, the ceiling is

$$\widehat{C}_{sc} = \left(\frac{1 + \hat{d}_{\text{short}}}{2} \right)^2 \geq C_0.$$

The archive's reported margin uses its own pooled factorized lower bound $\underline{s}_M^{\text{CP}}$:

$$\mathcal{M}_0 = \underline{s}_M^{\text{CP}} - C_0, \quad \mathcal{M}_{\text{sc}} = \underline{s}_M^{\text{CP}} - \widehat{C}_{\text{sc}}.$$

We use these reported boundaries only as a design-sensitivity analysis and do not numerically merge $\underline{s}_M^{\text{CP}}$ with the matched-separability archive's group-averaged lower bound. These margins do not identify the internal mechanism.

S2.6 Group Aggregation and Repeated-Success Outputs

Let $\mathcal{G}_c$ contain the groups with at least one complete matched pair. For each group g , let n_g be the number of complete pairs, S_g the number whose two labels are both correct, and $a_g = S_g/n_g$. The equal-group and equal-pair summaries are

$$\text{GADR} = \frac{1}{|\mathcal{G}_c|} \sum_{g \in \mathcal{G}_c} a_g, \quad \text{ADR} = \frac{\sum_g S_g}{\sum_g n_g}.$$

For $v_g = a_g(1 - a_g)$ and $\lambda \geq 0$, the analysis also exports

$$\text{SGADR}_\lambda = \frac{1}{|\mathcal{G}_c|} \sum_{g \in \mathcal{G}_c} (a_g - \lambda v_g), \quad \text{SGADR}_{\star, \lambda} = \text{ADR} - \lambda \text{ADR}(1 - \text{ADR}).$$

At $\lambda = 1$, these reduce to the macro-average of a_g^2 and ADR^2 , respectively. We additionally export between-group variance and

$$A_{\geq k} = \frac{1}{|\mathcal{G}_c|} \sum_{g \in \mathcal{G}_c} \mathbf{1}[S_g \geq k], \quad k = 1, \dots, 10.$$

The output bundle keeps the three domains separate and includes overall, model-size, per-group, lambda-sweep, completion, and applicability records. The primary files contain 19, 66, and 45 model-size rows for 3-SAT, 2-SAT, and graph coloring. The 3-SAT and graph-coloring recomputations match their archived summaries. For 2-SAT, the current predictions differ from the older export in 57 cells across four models. The domain-neutral implementation and the repository's original 2-SAT implementation agree on the current inputs, so the supplement retains the older differences as a versioning audit and reports the current-input values.

S3 Aggregate Secondary Results

S3.1 Selected Strict Ordering Reversals

S4 Threshold and Shortcut-Account Sensitivity

S4.1 Threshold-Parametric Matched-Separability Boundaries

S4.2 Ideal-Floor Versus Residual-Shortcut Boundaries

S4.3 Point Estimate Versus Confidence Requirement

S5 Complete Constraint Tables

Table 1. Metric level relative to the positive-side witness reference. Counts are computed row by row at source precision.

Domain	Metric	Rows	$M > W$	$M = W$	$M < W$	Mean ($M - W$)
3-SAT	Acc	19	18	1	0	0.322
3-SAT	ADR	19	15	4	0	0.077
3-SAT	$\hat{s}$	19	15	2	2	0.079
3-SAT	L_{CP}	19	10	0	9	-0.010
3-SAT	G-mean	19	18	1	0	0.234
3-SAT	BA	19	18	1	0	0.323
3-SAT	MCC	19	14	2	3	0.087
2-SAT	Acc	59	55	4	0	0.319
2-SAT	ADR	59	47	12	0	0.057
2-SAT	$\hat{s}$	59	47	12	0	0.057
2-SAT	L_{CP}	59	32	1	26	0.015
2-SAT	G-mean	59	54	5	0	0.205
2-SAT	BA	59	55	4	0	0.322
2-SAT	MCC	59	43	9	7	0.069
Graph coloring	Acc	45	45	0	0	0.490
Graph coloring	ADR	45	40	5	0	0.155
Graph coloring	$\hat{s}$	45	41	2	2	0.159
Graph coloring	L_{CP}	45	39	2	4	0.114
Graph coloring	G-mean	45	43	2	0	0.359
Graph coloring	BA	45	45	0	0	0.488
Graph coloring	MCC	43	35	2	6	0.134

Table 2. Strict pairwise model-ordering discordance with the witness reference. Each cell is discordant/strict-comparable pairs; source-precision ties are excluded.

Domain	Acc	F1 ₊	F1 ₋	ADR	$\hat{s}$	L_{CP}	G-mean	BA	MCC
3-SAT	3/16	4/16	2/16	4/16	4/16	3/16	4/16	3/16	3/16
2-SAT	10/115	25/116	10/111	11/116	12/116	9/116	13/116	12/116	14/116
Graph coloring	25/62	34/62	21/63	24/62	26/63	27/63	30/63	24/62	25/62

Table 3. Sensitivity to Clopper–Pearson versus Wilson lower bounds.

Domain	Mean abs. diff.	Median	95th pct.	Maximum
3-SAT	0.0043	0.0050	0.0073	0.010
2-SAT	0.0050	0.0060	0.0100	0.011
Graph coloring	0.0040	0.0030	0.0100	0.011

Table 4. Pair-rematching sensitivity, $\Delta_{\text{match}} = \text{ADR}_{\text{max}} - \text{ADR}_{\text{min}}$.

Domain	Rows	Mean	Median	Maximum	Nonzero rows
3-SAT	19	0.026	0.021	0.080	78.9%
2-SAT	59	0.023	0.000	0.182	42.4%
Graph coloring	45	0.072	0.019	0.364	73.3%

Table 5. Largest pair-rematching spreads in each domain.

Domain	Model	N	ADR_{pair}	ADR_{max}	ADR_{min}	Δ_{match}
3-SAT	DeepSeek-Reasoner	50	0.090	0.110	0.030	0.080
3-SAT	Opus 4.7 (med. think)	50	0.200	0.220	0.140	0.080
3-SAT	Opus 4.7 (med. think)	60	0.156	0.169	0.117	0.052
3-SAT	DeepSeek-Reasoner	60	0.086	0.086	0.038	0.048
3-SAT	DeepSeek-Reasoner	25	0.289	0.299	0.258	0.041
2-SAT	gpt-oss-20b (reason)	25	0.227	0.364	0.182	0.182
2-SAT	Qwen3-14B (think)	50	0.154	0.231	0.063	0.168
2-SAT	Qwen3-14B (think)	40	0.203	0.238	0.077	0.161
2-SAT	Qwen3-32B (think)	50	0.143	0.185	0.042	0.143
2-SAT	gpt-oss-20b (reason)	10	0.727	0.818	0.682	0.136
Graph coloring	Llama-Nemo-49B (r-on)	60	0.000	0.364	0.000	0.364
Graph coloring	Llama-Nemo-49B (r-on)	50	0.282	0.473	0.191	0.282
Graph coloring	QwQ-32B (think)	40	0.309	0.436	0.155	0.281
Graph coloring	Llama-Nemo-49B (r-on)	40	0.136	0.273	0.000	0.273
Graph coloring	Qwen3-32B (think)	20	0.278	0.389	0.141	0.248

Table 6. Representative model-order reversals relative to the witness reference. Arrows indicate the ordering induced by the displayed metric and by W .

Domain/size	Models	Output-metric ordering	Witness ordering
3-SAT, $N = 25$	Opus 4.7 vs. GPT-5	Acc 0.935 > 0.864; ADR 0.759 > 0.736; $\hat{s}$ 0.783 > 0.732	$W_{\text{Opus}} = 0.655 < W_{\text{GPT}} = 0.673$
Graph coloring, $N = 10$	Llama-Nemo-49B vs. GPT-OSS-20B	Acc 0.909 > 0.886; L_{CP} 0.751 > 0.710	$W_{\text{Llama}} = 0.227 < W_{\text{GPT-OSS}} = 0.818$
Graph coloring, $N = 8$	Llama-Nemo-49B vs. Qwen3-32B	Acc 0.932 > 0.824; ADR 0.864 > 0.647	$W_{\text{Llama}} = 0.136 < W_{\text{Qwen}} = 0.482$
2-SAT, $N = 30$	Qwen3-32B vs. Llama-Nemotron-Super-49B (reason)	ADR 0.157 < 0.250; $\hat{s}$ 0.156 < 0.182	$W_{\text{Qwen}} = 0.083 > W_{\text{Llama}} = 0.000$

Table 7. 3-SAT prefix boundaries $N^*(\tau)$ from the matched-separability archive.

Model	$N^*(0.50)$	$N^*(0.25)$	$N^*(0.10)$	$N^*(0.05)$
GPT-5	25	30	40	60
Opus 4.7 medium-thinking	25	25	50	60
DeepSeek-Reasoner	10	10	25	25

Table 8. 2-SAT prefix boundaries $N^*(\tau)$ from the matched-separability archive.

Model	$N^*(0.50)$	$N^*(0.25)$	$N^*(0.10)$	$N^*(0.05)$
Qwen3-14B thinking	15	20	50	50
Qwen3-32B thinking	10	15	20	30
QwQ-32B thinking	10	15	20	25

Table 9. Positive-margin boundaries from the shortcut-margin archive. $\mathcal{M}_0$ uses the ideal 0.25 floor; $\mathcal{M}_{sc}$ charges the measured residual structural shortcut ceiling. A dash means no positive boundary on the tested grid.

Task	Model	$\mathcal{M}_0$ boundary	$\mathcal{M}_{sc}$ boundary	Shift
3-SAT	GPT-5	30	25	−5
3-SAT	Opus 4.7 medium-thinking	25	25	0
3-SAT	DeepSeek-Reasoner	10	10	0
2-SAT	Qwen3-32B thinking	15	10	−5
2-SAT	gpt-oss-20b reasoning	20	15	−5
2-SAT	Nemotron-H-47B (r-on)	–	–	–
2-SAT	LN-Super-49B (r-on)	10	8	−2
Graph coloring	Qwen3-32B thinking	10	8	−2
Graph coloring	gpt-oss-20b reasoning	10	10	0
Graph coloring	Nemotron-H-47B (r-on)	15	10	−5
Graph coloring	LN-Super-49B (r-on)	10	10	0

Table 10. Examples in which an ideal-floor point margin is positive but its Clopper–Pearson lower-bound margin is negative.

Task	Model/size	Point $\mathcal{M}_0$	CP $\mathcal{M}_0$
3-SAT	GPT-5, $N = 40$	+0.053	−0.040
2-SAT	gpt-oss-20b, $N = 25$	+0.156	−0.045

Table 11. Complete 3-SAT model-size table from the matched-separability archive. W is the positive-side witness-aware rate: ADR^{+w} for SAT and ADR^{+a} for graph coloring. Values are reproduced at the three-decimal precision of the source export.

Model	N	Acc	F1 ₊	F1 ₋	ADR	W	$\hat{s}$	L_{CP}	L_W	BA	MCC	ADR _{pair}	ADR _{max}	ADR _{min}
GPT-5	5	1.000	1.000	1.000	1.000	1.000	1.000	0.929	0.927	1.000	1.000	1.000	1.000	1.000
GPT-5	10	0.990	0.991	0.989	0.989	0.989	0.981	0.897	0.897	0.991	0.980	0.989	0.989	0.989
GPT-5	25	0.864	0.847	0.877	0.736	0.673	0.732	0.612	0.615	0.864	0.745	0.736	0.755	0.727
GPT-5	30	0.735	0.651	0.787	0.486	0.440	0.480	0.367	0.372	0.734	0.533	0.486	0.495	0.468
GPT-5	40	0.645	0.466	0.735	0.309	0.164	0.303	0.210	0.216	0.645	0.393	0.309	0.309	0.291
GPT-5	50	0.573	0.288	0.695	0.173	0.064	0.168	0.099	0.105	0.573	0.242	0.173	0.173	0.145
GPT-5	60	0.555	0.222	0.688	0.127	0.000	0.125	0.067	0.072	0.555	0.210	0.127	0.127	0.118
Opus 4.7 (med. think)	5	0.945	0.947	0.942	0.897	0.813	0.893	0.782	0.784	0.945	0.892	0.897	0.907	0.888
Opus 4.7 (med. think)	10	0.959	0.960	0.958	0.917	0.861	0.919	0.816	0.817	0.959	0.918	0.917	0.926	0.917
Opus 4.7 (med. think)	25	0.935	0.956	0.875	0.759	0.655	0.783	0.592	0.602	0.894	0.838	0.759	0.759	0.759
Opus 4.7 (med. think)	50	0.527	0.620	0.372	0.200	0.060	0.230	0.124	0.131	0.555	0.134	0.200	0.220	0.140
Opus 4.7 (med. think)	60	0.530	0.656	0.256	0.156	0.000	0.152	0.078	0.084	0.549	0.157	0.156	0.169	0.117
DeepSeek-Reasoner	5	0.986	0.986	0.987	0.973	0.973	0.973	0.892	0.892	0.986	0.973	0.973	0.973	0.973
DeepSeek-Reasoner	10	0.876	0.865	0.886	0.780	0.780	0.761	0.641	0.644	0.879	0.774	0.780	0.780	0.770
DeepSeek-Reasoner	25	0.626	0.454	0.716	0.289	0.144	0.290	0.197	0.203	0.633	0.349	0.289	0.299	0.258
DeepSeek-Reasoner	30	0.515	0.154	0.660	0.073	0.010	0.082	0.036	0.041	0.519	0.075	0.073	0.073	0.052
DeepSeek-Reasoner	40	0.524	0.183	0.664	0.093	0.031	0.104	0.048	0.053	0.520	0.072	0.093	0.103	0.082
DeepSeek-Reasoner	50	0.486	0.194	0.622	0.090	0.000	0.102	0.052	0.057	0.496	-0.011	0.090	0.110	0.030
DeepSeek-Reasoner	60	0.509	0.160	0.653	0.086	0.010	0.086	0.039	0.044	0.513	0.049	0.086	0.086	0.038

Table 12. Complete 2-SAT model-size table from the matched-separability archive. W is the positive-side witness-aware rate: ADR^{+w} for SAT and ADR^{+a} for graph coloring. Values are reproduced at the three-decimal precision of the source export.

Model	N	Acc	F1 ₊	F1 ₋	ADR	W	$\hat{s}$	L_{CP}	L_W	BA	MCC	ADR _{pair}	ADR _{max}	ADR _{min}
gpt-oss-20b (reason)	3	1.000	1.000	1.000	1.000	1.000	1.000	0.976	0.969	1.000	1.000	1.000	1.000	1.000
gpt-oss-20b (reason)	5	1.000	1.000	1.000	1.000	1.000	1.000	0.976	0.969	1.000	1.000	1.000	1.000	1.000
gpt-oss-20b (reason)	8	0.955	0.952	0.957	0.909	0.864	0.909	0.863	0.855	0.955	0.913	0.909	0.909	0.909
gpt-oss-20b (reason)	10	0.841	0.844	0.837	0.727	0.682	0.707	0.631	0.620	0.841	0.683	0.727	0.818	0.682
gpt-oss-20b (reason)	15	0.773	0.737	0.800	0.545	0.500	0.562	0.485	0.475	0.773	0.567	0.545	0.636	0.545
gpt-oss-20b (reason)	20	0.773	0.706	0.815	0.545	0.409	0.545	0.467	0.457	0.773	0.612	0.545	0.545	0.545
gpt-oss-20b (reason)	25	0.568	0.457	0.642	0.227	0.182	0.281	0.216	0.210	0.568	0.149	0.227	0.364	0.182
gpt-oss-20b (reason)	30	0.568	0.296	0.689	0.136	0.045	0.169	0.118	0.116	0.568	0.215	0.136	0.182	0.136
gpt-oss-20b (reason)	40	0.545	0.286	0.667	0.182	0.000	0.165	0.112	0.110	0.545	0.132	0.182	0.182	0.091
gpt-oss-20b (reason)	50	0.500	0.083	0.656	0.045	0.000	0.045	0.026	0.026	0.500	0.000	0.045	0.045	0.045
Qwen3-14B (think)	3	0.990	0.990	0.990	0.981	0.981	0.981	0.950	0.943	0.990	0.981	0.981	0.981	0.981
Qwen3-14B (think)	5	0.974	0.973	0.975	0.948	0.942	0.948	0.909	0.901	0.974	0.949	0.948	0.948	0.948
Qwen3-14B (think)	8	0.974	0.973	0.975	0.948	0.942	0.948	0.907	0.899	0.974	0.949	0.948	0.948	0.948
Qwen3-14B (think)	10	0.942	0.938	0.945	0.883	0.877	0.883	0.833	0.824	0.942	0.889	0.883	0.883	0.883
Qwen3-14B (think)	15	0.844	0.815	0.865	0.688	0.656	0.688	0.620	0.611	0.844	0.724	0.688	0.688	0.688
Qwen3-14B (think)	20	0.669	0.519	0.748	0.338	0.286	0.348	0.282	0.276	0.669	0.432	0.338	0.357	0.338
Qwen3-14B (think)	25	0.604	0.371	0.711	0.214	0.156	0.230	0.179	0.175	0.604	0.309	0.214	0.234	0.208
Qwen3-14B (think)	30	0.516	0.280	0.636	0.162	0.058	0.153	0.105	0.103	0.516	0.043	0.162	0.188	0.078
Qwen3-14B (think)	40	0.490	0.338	0.585	0.203	0.000	0.184	0.136	0.133	0.495	-0.012	0.203	0.238	0.077
Qwen3-14B (think)	50	0.479	0.313	0.580	0.154	0.000	0.163	0.118	0.115	0.479	-0.048	0.154	0.231	0.063
Qwen3-32B (think)	3	0.979	0.979	0.980	0.959	0.959	0.959	0.922	0.914	0.979	0.959	0.959	0.959	0.959
Qwen3-32B (think)	5	0.955	0.952	0.957	0.909	0.909	0.909	0.864	0.855	0.955	0.913	0.909	0.909	0.909
Qwen3-32B (think)	8	0.926	0.920	0.931	0.851	0.843	0.851	0.798	0.790	0.926	0.861	0.851	0.851	0.851
Qwen3-32B (think)	10	0.905	0.895	0.913	0.810	0.777	0.810	0.753	0.744	0.905	0.825	0.810	0.810	0.810
Qwen3-32B (think)	15	0.748	0.663	0.799	0.496	0.430	0.496	0.422	0.413	0.748	0.574	0.496	0.496	0.496
Qwen3-32B (think)	20	0.624	0.397	0.727	0.248	0.207	0.248	0.189	0.185	0.624	0.376	0.248	0.248	0.248
Qwen3-32B (think)	25	0.568	0.246	0.698	0.142	0.092	0.140	0.098	0.096	0.567	0.254	0.142	0.142	0.133
Qwen3-32B (think)	30	0.574	0.270	0.700	0.157	0.083	0.156	0.112	0.110	0.574	0.270	0.157	0.157	0.149
Qwen3-32B (think)	40	0.525	0.173	0.667	0.099	0.000	0.096	0.067	0.066	0.525	0.094	0.099	0.099	0.083
Qwen3-32B (think)	50	0.467	0.273	0.579	0.143	0.000	0.129	0.089	0.087	0.464	-0.083	0.143	0.185	0.042
QwQ-32B (think)	3	1.000	1.000	1.000	1.000	1.000	1.000	0.976	0.969	1.000	1.000	1.000	1.000	1.000
QwQ-32B (think)	5	0.990	0.990	0.990	0.980	0.980	0.980	0.949	0.942	0.990	0.980	0.980	0.980	0.980
QwQ-32B (think)	8	0.980	0.980	0.980	0.960	0.960	0.960	0.924	0.916	0.980	0.960	0.960	0.960	0.960
QwQ-32B (think)	10	0.949	0.947	0.952	0.899	0.869	0.899	0.857	0.849	0.949	0.904	0.899	0.899	0.899
QwQ-32B (think)	15	0.798	0.747	0.832	0.596	0.545	0.596	0.521	0.511	0.798	0.652	0.596	0.596	0.596
QwQ-32B (think)	20	0.641	0.441	0.736	0.283	0.242	0.283	0.222	0.218	0.641	0.406	0.283	0.283	0.283
QwQ-32B (think)	25	0.576	0.276	0.700	0.162	0.111	0.161	0.122	0.119	0.576	0.270	0.162	0.162	0.152
QwQ-32B (think)	30	0.515	0.059	0.673	0.030	0.000	0.030	0.017	0.017	0.515	0.124	0.030	0.030	0.030
QwQ-32B (think)	40	0.487	0.093	0.642	0.033	0.011	0.039	0.018	0.019	0.509	0.043	0.033	0.044	0.022
QwQ-32B (think)	50	0.506	0.044	0.667	0.023	0.000	0.023	0.013	0.013	0.506	0.044	0.023	0.023	0.023

Continued on next page

Table 12 continued														
Model	N	Acc	F1 ₊	F1 _−	ADR	W	$\hat{s}$	L_{CP}	L_W	BA	MCC	ADR _{pair}	ADR _{max}	ADR _{min}
Llama-Nemotron-Super-49B (reason)	3	1.000	1.000	1.000	1.000	0.909	1.000	0.976	0.969	1.000	1.000	1.000	1.000	1.000
Llama-Nemotron-Super-49B (reason)	5	1.000	1.000	1.000	1.000	1.000	1.000	0.976	0.969	1.000	1.000	1.000	1.000	1.000
Llama-Nemotron-Super-49B (reason)	8	0.864	0.842	0.880	0.727	0.727	0.727	0.653	0.642	0.864	0.756	0.727	0.727	0.727
Llama-Nemotron-Super-49B (reason)	10	0.773	0.706	0.815	0.545	0.545	0.545	0.467	0.457	0.773	0.612	0.545	0.545	0.545
Llama-Nemotron-Super-49B (reason)	15	0.727	0.625	0.786	0.455	0.364	0.455	0.377	0.369	0.727	0.542	0.455	0.455	0.455
Llama-Nemotron-Super-49B (reason)	20	0.818	0.778	0.846	0.636	0.545	0.636	0.558	0.548	0.818	0.683	0.636	0.636	0.636
Llama-Nemotron-Super-49B (reason)	25	0.773	0.706	0.815	0.545	0.182	0.545	0.467	0.457	0.773	0.612	0.545	0.545	0.545
Llama-Nemotron-Super-49B (reason)	30	0.400	0.308	0.471	0.250	0.000	0.182	0.094	0.095	0.591	0.237	0.250	0.250	0.250
Nemotron-H-47B-Reasoning (reason)	3	0.524	0.657	0.224	0.077	0.046	0.080	0.069	0.068	0.524	0.076	0.077	0.087	0.074
Nemotron-H-47B-Reasoning (reason)	5	0.510	0.668	0.066	0.023	0.004	0.027	0.019	0.019	0.510	0.062	0.023	0.031	0.021
Nemotron-H-47B-Reasoning (reason)	8	0.499	0.666	−	0.000	0.000	0.000	0.000	0.000	0.499	−0.031	0.000	0.000	0.000
Nemotron-H-47B-Reasoning (reason)	10	0.497	0.664	0.004	0.002	0.000	0.002	0.001	0.001	0.497	−0.042	0.002	0.002	0.002
Nemotron-H-47B-Reasoning (reason)	15	0.513	0.668	0.087	0.027	0.000	0.030	0.024	0.024	0.513	0.070	0.027	0.031	0.025
Nemotron-H-47B-Reasoning (reason)	20	0.510	0.664	0.093	0.043	0.000	0.043	0.037	0.036	0.510	0.049	0.043	0.046	0.041
Nemotron-H-47B-Reasoning (reason)	25	0.526	0.666	0.183	0.070	0.000	0.071	0.058	0.057	0.526	0.096	0.070	0.079	0.062
Nemotron-H-47B-Reasoning (reason)	30	0.503	0.665	0.034	0.015	0.000	0.017	0.012	0.012	0.503	0.024	0.015	0.017	0.015
Nemotron-H-47B-Reasoning (reason)	40	0.527	0.673	0.147	0.074	0.000	0.074	0.055	0.054	0.527	0.120	0.074	0.081	0.062
Nemotron-H-47B-Reasoning (reason)	50	0.516	0.667	0.120	0.046	0.000	0.052	0.042	0.041	0.516	0.076	0.046	0.062	0.041
Nemotron-H-47B-Reasoning (reason)	60	0.537	0.674	0.206	0.105	0.000	0.107	0.085	0.083	0.535	0.125	0.105	0.119	0.093

Table 13. Complete graph-coloring model–size table from the matched-separability archive. W is the positive-side witness-aware rate: ADR^{+w} for SAT and ADR^{+a} for graph coloring. Values are reproduced at the three-decimal precision of the source export.

Model	N	Acc	$F1_+$	$F1_-$	ADR	W	$\hat{s}$	L_{CP}	L_W	BA	MCC	ADR_{pair}	ADR_{max}	ADR_{min}
GPT-OSS-20B (think)	8	0.953	0.952	0.955	0.905	0.857	0.909	0.852	0.843	0.955	0.911	0.905	0.905	0.905
GPT-OSS-20B (think)	10	0.886	0.878	0.894	0.818	0.781	0.710	0.700	0.886	0.780	0.818	0.818	0.818	0.773
GPT-OSS-20B (think)	15	0.545	0.286	0.667	0.182	0.136	0.165	0.130	0.127	0.545	0.132	0.182	0.182	0.136
GPT-OSS-20B (think)	20	0.659	0.462	0.750	0.316	0.316	0.313	0.242	0.236	0.650	0.424	0.322	0.322	0.322
GPT-OSS-20B (think)	30	0.528	–	0.691	0.000	0.000	0.000	0.000	0.000	0.500	–	0.000	0.000	0.000
GPT-OSS-20B (think)	40	0.571	0.118	0.717	0.071	0.000	0.062	0.036	0.036	0.531	0.187	0.062	0.062	0.062
GPT-OSS-20B (think)	50	0.440	0.222	0.562	0.182	0.000	0.149	0.098	0.096	0.481	-0.053	0.182	0.182	0.000
GPT-OSS-20B (think)	60	0.550	0.308	0.667	0.111	0.000	0.180	0.120	0.117	0.550	0.140	0.111	0.222	0.111
Llama-Nemo-49B (r-on)	8	0.932	0.927	0.936	0.864	0.136	0.864	0.802	0.792	0.932	0.872	0.864	0.864	0.864
Llama-Nemo-49B (r-on)	10	0.909	0.900	0.917	0.818	0.227	0.818	0.751	0.740	0.909	0.832	0.818	0.818	0.818
Llama-Nemo-49B (r-on)	15	0.545	0.231	0.677	0.136	0.091	0.124	0.092	0.090	0.545	0.158	0.136	0.136	0.091
Llama-Nemo-49B (r-on)	20	0.545	0.167	0.688	0.091	0.045	0.091	0.063	0.062	0.545	0.218	0.091	0.091	0.091
Llama-Nemo-49B (r-on)	30	0.545	0.444	0.615	0.182	0.000	0.264	0.209	0.205	0.545	0.098	0.182	0.364	0.182
Llama-Nemo-49B (r-on)	40	0.364	0.364	0.364	0.136	0.000	0.132	0.087	0.085	0.364	-0.273	0.136	0.273	0.000
Llama-Nemo-49B (r-on)	50	0.632	0.682	0.562	0.312	0.000	0.337	0.249	0.243	0.622	0.244	0.282	0.473	0.191
Llama-Nemo-49B (r-on)	60	0.409	0.381	0.435	0.000	0.000	0.165	0.112	0.110	0.409	-0.183	0.000	0.364	0.000
Nemo-H-47B (r-on)	8	0.668	0.594	0.720	0.338	0.002	0.346	0.293	0.288	0.668	0.362	0.338	0.356	0.337
Nemo-H-47B (r-on)	10	0.641	0.638	0.645	0.301	0.001	0.302	0.251	0.246	0.641	0.283	0.301	0.321	0.286
Nemo-H-47B (r-on)	15	0.582	0.482	0.650	0.193	0.000	0.202	0.160	0.157	0.583	0.179	0.193	0.227	0.171
Nemo-H-47B (r-on)	20	0.556	0.406	0.646	0.144	0.000	0.153	0.121	0.118	0.556	0.131	0.144	0.173	0.130
Nemo-H-47B (r-on)	30	0.537	0.346	0.641	0.116	0.000	0.126	0.097	0.095	0.537	0.090	0.116	0.152	0.093
Nemo-H-47B (r-on)	40	0.505	0.067	0.663	0.021	0.000	0.022	0.017	0.016	0.505	0.029	0.021	0.028	0.017
Nemo-H-47B (r-on)	50	0.503	0.030	0.666	0.009	0.000	0.009	0.007	0.007	0.503	0.027	0.009	0.010	0.008
Nemo-H-47B (r-on)	60	0.500	–	0.667	0.000	0.000	0.000	0.000	0.000	0.500	–	0.000	0.000	0.000
Qwen3-14B (think)	8	0.853	0.829	0.871	0.706	0.473	0.708	0.638	0.629	0.853	0.737	0.706	0.711	0.706
Qwen3-14B (think)	10	0.755	0.680	0.802	0.517	0.351	0.514	0.445	0.436	0.757	0.585	0.515	0.518	0.515
Qwen3-14B (think)	15	0.568	0.478	0.632	0.311	0.096	0.297	0.233	0.227	0.570	0.150	0.311	0.385	0.177
Qwen3-14B (think)	20	0.516	0.604	0.379	0.225	0.000	0.212	0.157	0.153	0.516	0.037	0.225	0.290	0.093
Qwen3-14B (think)	30	0.501	0.660	0.066	0.030	0.000	0.034	0.020	0.020	0.501	0.007	0.030	0.035	0.028
Qwen3-14B (think)	40	0.500	0.657	0.079	0.035	0.000	0.041	0.023	0.023	0.500	0.000	0.035	0.043	0.023
Qwen3-14B (think)	50	0.539	0.693	0.072	0.039	0.000	0.038	0.021	0.021	0.503	0.015	0.039	0.039	0.026
Qwen3-32B (think)	8	0.824	0.788	0.850	0.647	0.482	0.648	0.576	0.566	0.823	0.687	0.645	0.655	0.645
Qwen3-32B (think)	10	0.692	0.564	0.762	0.389	0.308	0.393	0.325	0.318	0.692	0.474	0.389	0.399	0.384
Qwen3-32B (think)	15	0.580	0.370	0.685	0.225	0.096	0.224	0.175	0.171	0.581	0.220	0.223	0.245	0.173
Qwen3-32B (think)	20	0.548	0.475	0.604	0.284	0.046	0.272	0.210	0.205	0.550	0.104	0.278	0.389	0.141
Qwen3-32B (think)	30	0.497	0.586	0.361	0.214	0.000	0.199	0.144	0.141	0.497	-0.006	0.214	0.283	0.059
Qwen3-32B (think)	40	0.505	0.667	0.041	0.021	0.000	0.021	0.012	0.012	0.505	0.043	0.021	0.021	0.016
Qwen3-32B (think)	50	0.583	0.722	0.167	0.091	0.000	0.091	0.052	0.052	0.545	0.227	0.091	0.091	0.091
QwQ-32B (think)	8	0.919	0.912	0.925	0.837	0.407	0.837	0.784	0.775	0.919	0.849	0.837	0.837	0.837
QwQ-32B (think)	10	0.831	0.796	0.855	0.661	0.331	0.661	0.588	0.578	0.831	0.703	0.661	0.661	0.661

Continued on next page

431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464

Table 13 continued														
Model	N	Acc	$F1_+$	$F1_-$	ADR	W	$\hat{s}$	L_{CP}	L_W	BA	MCC	ADR_{pair}	ADR_{max}	ADR_{min}
QwQ-32B (think)	15	0.598	0.339	0.711	0.206	0.115	0.204	0.155	0.151	0.598	0.316	0.206	0.206	0.196
QwQ-32B (think)	20	0.566	0.234	0.697	0.132	0.041	0.132	0.089	0.087	0.566	0.266	0.132	0.132	0.132
QwQ-32B (think)	30	0.512	0.224	0.645	0.124	0.008	0.113	0.080	0.078	0.512	0.037	0.124	0.140	0.050
QwQ-32B (think)	40	0.559	0.554	0.565	0.309	0.000	0.298	0.233	0.228	0.560	0.121	0.309	0.436	0.155
QwQ-32B (think)	50	0.500	0.645	0.154	0.091	0.000	0.083	0.046	0.046	0.500	0.000	0.091	0.091	0.000

S6 Architecture-Level Instantiations and Intervention Diagnostics

The main results of this paper are architecture independent. They apply whenever model behavior is evaluated through binary predictions or scalar scores. This section illustrates how the same limitations arise in three familiar model families: additive classifiers, attention-based models, and graph neural networks (GNNs). These examples do not establish additional impossibility results. Their purpose is to show why architectural complexity does not make a higher prediction metric value a reliable indicator of stronger reasoning.

Across these model families, semantic and non-semantic information can both affect the final score. If non-semantic information is more predictive on the benchmark distribution, a model that exploits it can outperform a model that attempts to capture the intended semantic relationship but does so imperfectly. The resulting metric ordering is valid as an ordering of benchmark performance, but it need not be an ordering of reasoning ability.

S6.1 Additive Classifiers over Feature Blocks

We first consider an additive classifier because the contributions of semantic and non-semantic information can be represented explicitly. This is a grouped abstraction of familiar linear, logistic, and additive models [1–3, 9].

Following the semantic/non-semantic decomposition in the main paper, let z denote the semantic factor that should determine the label, and let c denote non-semantic information that may also be available to the model. Let

$$a_z = \phi_z(z), \quad a_c = \phi_c(c)$$

be feature vectors derived from these components. Consider a model f with score

$$s_f(x) = w_z^\top a_z + w_c^\top a_c + b$$

and binary prediction

$$\hat{y} = 1[s_f(x) \geq 0].$$

Define the two score contributions by

$$S_z(x) = w_z^\top a_z, \quad S_c(x) = w_c^\top a_c.$$

Then

$$s_f(x) = S_z(x) + S_c(x) + b.$$

This representation makes two earlier results concrete. First, distinct semantic and non-semantic contributions can yield the same scalar score and therefore the same prediction and ranking. This is an instance of the non-identifiability established in Proposition 1. Second, a larger non-semantic contribution may improve benchmark predictions or rankings even when it does not improve semantic reasoning.

For example, consider a semantic classifier and a shortcut classifier with class-conditional error rates

$$(\delta_1, \delta_2) \quad \text{and} \quad (\epsilon_1, \epsilon_2),$$

respectively. As shown in Proposition 2, the shortcut classifier receives higher accuracy whenever

$$\pi\epsilon_1 + (1 - \pi)\epsilon_2 < \pi\delta_1 + (1 - \pi)\delta_2,$$

where $\pi = \Pr[Y = 1]$. Analogous comparisons hold for balanced accuracy, G-mean, Youden's J , F1, and other prediction metrics.

The additive model makes the interpretation especially clear. The shortcut classifier is ranked higher because its non-semantic features produce a more favorable benchmark error profile. The

higher metric value does not imply that the classifier has captured the semantic relationship that determines the labels.

A similar conclusion applies to score-based metrics. If the non-semantic contribution ranks positive and negative benchmark instances more reliably than the semantic contribution, it can increase AUROC. Thus, observing the unthresholded score avoids one source of information loss, but it does not make metric ordering monotone with reasoning ability.

S6.2 Attention-Based Models

Attention-based models provide a less explicit version of the same phenomenon. Let an input be

$$x = (t_1, \dots, t_n),$$

where t_j is the token occurrence at position j . A Transformer maps the input to hidden representations through repeated attention and feed-forward transformations [7]. Abstractly,

$$H^{(0)} = \text{Embed}(x),$$

$$H^{(\ell+1)} = \text{Block}_\ell(H^{(\ell)}), \quad \ell = 0, \dots, L-1.$$

After the final layer, the model constructs an instance representation

$$h_x = \text{Pool}(H^{(L)})$$

and a scalar score

$$s_f(x) = w_{\text{cls}}^\top h_x + b_{\text{cls}}.$$

In software-engineering tasks, some input information bears directly on the intended semantics, while other information may correlate with the labels without determining them. In vulnerability detection, for example, relevant semantic information may include relationships among sources, checks, sanitizers, and sinks. Non-semantic information may include identifier names, comments, API names, formatting conventions, or dataset-construction artifacts. In SAT solving, semantic information is distributed across variables, literals, clauses, and connectives, while generator-specific formatting or structural regularities may provide non-semantic cues.

Attention-based models can use either kind of information. A model that captures the relevant semantic relationships may still make errors because those relationships are difficult to learn. Another model may exploit easier non-semantic correlations and obtain a better benchmark error profile or ranking. In that case,

$$M(B, f_c) > M(B, f_z)$$

may hold even though f_z relies more strongly on the intended semantic relationship.

This is the practically important implication of the attention instantiation. A higher accuracy, F1, ADR, or AUROC does not establish that the higher-scoring attention model reasons more effectively. It establishes only that the model performs better according to the observable criterion of the metric on the evaluated distribution.

Attention weights, gradients, and attribution methods such as Integrated Gradients, SHAP, and DeepLIFT can provide additional diagnostic information [4–6]. However, these analyses should not be treated as direct measures of reasoning. Their results depend on the explanation method, the selected representation, and how interactions among tokens are allocated. They can motivate further tests, but they do not by themselves justify interpreting a higher prediction metric as stronger reasoning.

S6.3 Graph Neural Networks and Message Passing

Graph neural networks instantiate the same limitation in models that operate over structured program representations. Let

$$G = (V, E, X)$$

be a program graph, where V is the node set, E is the edge set, and X_v is the initial feature vector of node v . A generic message-passing layer computes

$$m_v^{(k)} = \text{Agg}_{u \in N(v)} \text{Msg}_k(h_v^{(k)}, h_u^{(k)}, e_{uv}),$$

followed by

$$h_v^{(k+1)} = \text{Upd}_k(h_v^{(k)}, m_v^{(k)}).$$

After K layers, the model produces a graph representation

$$h_G = \text{Readout}_{v \in V} h_v^{(K)}$$

and score

$$s_f(G) = w^\top h_G + b.$$

Program graphs may encode AST, control-flow, data-flow, call, or dependence relationships. The semantic condition underlying a vulnerability may depend on a long-range relation, such as an untrusted value reaching a sink without sanitization, a missing bounds check before a memory access, or a particular alias being dereferenced under an unsafe path. These relationships can be difficult to represent and learn.

At the same time, graph inputs may contain easier non-semantic correlations, including risky API names, identifier features, project-specific graph patterns, or local motifs introduced by data extraction. A GNN that exploits these correlations may obtain a higher metric value than a detector that attempts to capture the relevant source-to-sink or control-flow relationship but does so imperfectly.

The resulting ranking is again non-monotone with respect to reasoning ability. A higher metric value can indicate that local or project-specific features are more predictive on the benchmark. It does not necessarily indicate that the model has captured the vulnerability condition represented by the graph.

Graph explanation methods can help diagnose this possibility. For example, GNNExplainer identifies a subgraph and feature subset that approximately preserve a model's prediction [8]. An explanation centered on a vulnerability-relevant slice may provide encouraging evidence, whereas one centered on unrelated names or local graph motifs may reveal a potential shortcut. Nevertheless, an explanation describes what influenced the model's output under the assumptions of the explanation method. It does not establish that the model's metric advantage reflects stronger vulnerability reasoning.

S6.4 Intervention-Based Comparison

The model-level examples suggest a common way to strengthen reasoning-oriented evaluation. Rather than interpreting an aggregate metric value in isolation, an evaluator can examine how model outputs respond to controlled semantic and non-semantic changes.

Let $x_{\Delta z}$ be an instance in which the semantic condition is changed while non-semantic information is held as fixed as possible. Let $x_{\Delta c}$ preserve the semantic condition while changing non-semantic information. Define the score changes

$$\Delta_z(x) = |s_f(x) - s_f(x_{\Delta z})|$$

and

$$\Delta_c(x) = |s_f(x) - s_f(x_{\Delta c})|.$$

A model whose outputs track the intended task semantics should generally respond to semantic changes that alter the correct answer and remain comparatively stable under semantics-preserving non-semantic changes. For binary predictions, analogous quantities can be defined through prediction flips:

$$F_z(x) = \mathbf{1}[f(x) \neq f(x_{\Delta z})],$$

$$F_c(x) = \mathbf{1}[f(x) \neq f(x_{\Delta c})].$$

Matched SAT/UNSAT formulas and vulnerable/fixed program variants approximate semantic interventions. Identifier renaming, formatting changes, irrelevant code transformations, and generator perturbations approximate semantics-preserving non-semantic interventions. Pair, group, and cluster metrics summarize success over such controlled instances.

These interventions provide a more informative evaluation design, but the limitations of controlled metrics established in the main paper still apply. A higher controlled score indicates stronger performance under the selected interventions and aggregation rule. It does not automatically imply stronger reasoning, because residual non-semantic correlations, intervention quality, dependence among instances, and threshold selection can still affect the ordering.

S6.5 Summary

The three instantiations differ in how information is represented and combined, but they support the same interpretation. Additive classifiers may combine explicit feature-block contributions. Attention-based models mix information across token representations. GNNs aggregate node and edge information through message passing. These architectural differences can help explain why a model succeeds, but they do not change what a prediction-output metric establishes.

Across all three settings, a higher metric value establishes better performance under the metric's observable benchmark criterion. It does not, without controlled interventions or additional task-specific evidence, establish stronger reasoning. Architecture-specific explanations and attributions can provide useful diagnostics, while witnesses and controlled semantic and non-semantic interventions provide stronger behavioral evidence. None of these observations makes an ordinary prediction metric a standalone measure of reasoning ability.

S7 Reproduction Notes and Remaining Requirements

The accompanying revision package includes CSV files transcribed from the source LaTeX tables and CSV outputs for the metric-level, ordering, pair-rematching, and confidence-bound analyses. The strict-ordering calculation should be rerun on full-precision values before submission. Per-instance predictions should then be used to verify every source-table cell.

The source matched-separability protocol describes a matched-pair cluster bootstrap for dependence within seed families. The supplied summary files do not contain per-instance outcomes or bootstrap replicates, so this revision does not report invented bootstrap values. A final artifact should resample matched pairs within groups, recompute the group-averaged $\hat{s}$, and report a dependence-robust lower quantile beside the Clopper–Pearson and Wilson reference bounds.

Witness validation should also be strengthened on the negative side. SAT experiments should elicit and check a normalized UNSAT proof when feasible; graph-coloring experiments should use a checkable non-colorability certificate or a trusted solver trace. Until then, W remains a positive-side constructive reference rather than a complete semantic oracle.

References

[1] Rong-En Fan, Kai-Wei Chang, Cho-Jui Hsieh, Xiang-Rui Wang, and Chih-Jen Lin. 2008. LIBLINEAR: A library for large linear classification. *the Journal of machine Learning research* 9 (2008), 1871–1874.

[2] Trevor J Hastie. 2017. Generalized additive models. *Statistical models in S* (2017), 249–307.

[3] Thorsten Joachims. 1998. Text categorization with support vector machines: Learning with many relevant features. In *European conference on machine learning*. Springer, 137–142.

[4] Scott M Lundberg and Su-In Lee. 2017. A unified approach to interpreting model predictions. *Advances in neural information processing systems* 30 (2017).

[5] Avanti Shrikumar, Peyton Greenside, and Anshul Kundaje. 2017. Learning important features through propagating activation differences. In *International conference on machine learning*. PMLR, 3145–3153.

[6] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. 2017. Axiomatic attribution for deep networks. In *International conference on machine learning*. PMLR, 3319–3328.

[7] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems* 30 (2017).

[8] Zhitao Ying, Dylan Bourgeois, Jiaxuan You, Marinka Zitnik, and Jure Leskovec. 2019. Gnnexplainer: Generating explanations for graph neural networks. *Advances in neural information processing systems* 32 (2019).

[9] Ming Yuan and Yi Lin. 2006. Model selection and estimation in regression with grouped variables. *Journal of the Royal Statistical Society Series B: Statistical Methodology* 68, 1 (2006), 49–67.